

Malware Classification Using Ember features and Feed Forward Neural Network

used a simple 3-layer feed-forward network with dropout and batch-norm.

At the end it can be seen that the training accuracy is 99% and validation accuracy is 92% , and testing accuracy is 92% The model is converging in the initial 5 epochs itself. Since a malware sample can belongs to multiple classes at the same time. and the label with max probability is considered as true label.the validation and test accuracy around 92% is still good. I belive that if we consider top 3 labels then the accuracy can be much improved.

the corressponding confusion matrix and classification reports and plots can be seen at the end.

the architecure implemented here is very simple and more similar to the one in this link:
<https://towardsdatascience.com/pytorch-tabular-multiclass-classification-9f8211a123ab>

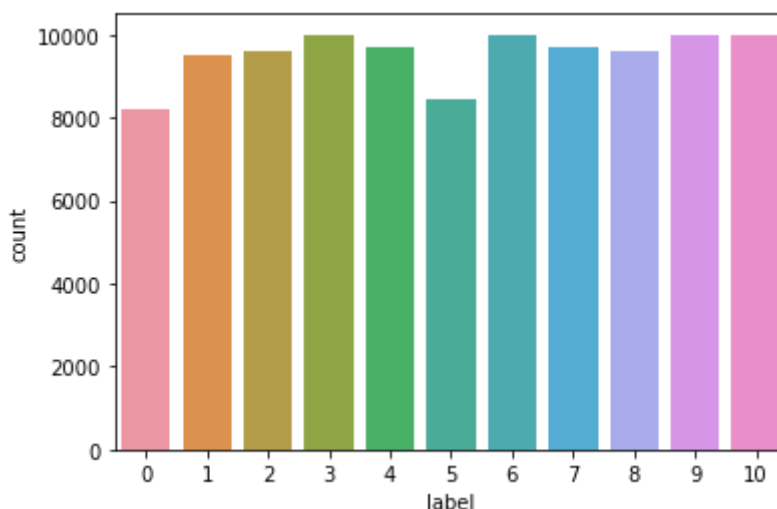
```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
from tqdm.notebook import tqdm
import matplotlib.pyplot as plt
import os
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, WeightedRandomSampler

from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

```
In [2]: sorel_dir = '../Dataset/sorel'
data = np.load(os.path.join(sorel_dir, 'sorel_data.npz'), allow_pickle=True)
data['arr_0'].shape, data['arr_1'].shape
```

```
Out[2]: ((104746, 2381), (104746, 13))
```

```
In [3]: df = pd.DataFrame(data=data['arr_1'][:, -1], columns=["label"])
sns.countplot(x = 'label', data=df);
```



```
In [4]: # Split into train+val and test
X_trainval, X_test, y_trainval, y_test = train_test_split(data['arr_0'], data['arr_1']

# Split train into train-val
X_train, X_val, y_train, y_val = train_test_split(X_trainval, y_trainval, test_size=
```

```
In [5]: scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_val = scaler.transform(X_val)
X_test = scaler.transform(X_test)
X_train, y_train = np.array(X_train), np.array(y_train)
X_val, y_val = np.array(X_val), np.array(y_val)
X_test, y_test = np.array(X_test), np.array(y_test)
```

```
In [7]: class ClassifierDataset(Dataset):
    def __init__(self, X_data, y_data):
        self.X_data = X_data
        self.y_data = y_data
    def __getitem__(self, index):
        return self.X_data[index], self.y_data[index]
    def __len__(self):
        return len(self.X_data)

class MulticlassClassification(nn.Module):
    def __init__(self, num_feature, num_class):
        super(MulticlassClassification, self).__init__()

        self.layer_1 = nn.Linear(num_feature, 512)
        self.layer_2 = nn.Linear(512, 128)
        self.layer_3 = nn.Linear(128, 64)
        self.layer_out = nn.Linear(64, num_class)

        self.relu = nn.ReLU()
        self.dropout = nn.Dropout(p=0.2)
        self.batchnorm1 = nn.BatchNorm1d(512)
        self.batchnorm2 = nn.BatchNorm1d(128)
        self.batchnorm3 = nn.BatchNorm1d(64)

    def forward(self, x):
        x = self.layer_1(x)
        x = self.batchnorm1(x)
        x = self.relu(x)

        x = self.layer_2(x)
        x = self.batchnorm2(x)
        x = self.relu(x)
        x = self.dropout(x)

        x = self.layer_3(x)
        x = self.batchnorm3(x)
        x = self.relu(x)
        x = self.dropout(x)

        x = self.layer_out(x)

        return x
```

```
In [8]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
values, counts = np.unique(y_train, return_counts=True)
class_weights = 1./torch.tensor(counts, dtype=torch.float)
print(class_weights)

tensor([0.0002, 0.0001, 0.0001, 0.0001, 0.0001, 0.0002, 0.0001, 0.0001, 0.0001,
        0.0001, 0.0001])
```

```
In [9]: EPOCHS = 50
        BATCH_SIZE = 16
        LEARNING_RATE = 0.0007
```

```
NUM_FEATURES = 2381
NUM_CLASSES = 11
```

```
In [10]: train_dataset = ClassifierDataset(torch.from_numpy(X_train).float(), torch.from_numpy(y_train).float())
val_dataset = ClassifierDataset(torch.from_numpy(X_val).float(), torch.from_numpy(y_val).float())
test_dataset = ClassifierDataset(torch.from_numpy(X_test).float(), torch.from_numpy(y_test).float())
train_loader = DataLoader(dataset=train_dataset,
                           batch_size=BATCH_SIZE,
                           shuffle=True)
val_loader = DataLoader(dataset=val_dataset, batch_size=1)
test_loader = DataLoader(dataset=test_dataset, batch_size=1)
```

```
In [31]: model = MulticlassClassification(num_feature = NUM_FEATURES, num_class=NUM_CLASSES)
model.load_state_dict(torch.load(os.path.join(sorel_dir, 'classification.pth')))
model.eval()
model.to(device)
```

```
Out[31]: MulticlassClassification(
  (layer_1): Linear(in_features=2381, out_features=512, bias=True)
  (layer_2): Linear(in_features=512, out_features=128, bias=True)
  (layer_3): Linear(in_features=128, out_features=64, bias=True)
  (layer_out): Linear(in_features=64, out_features=11, bias=True)
  (relu): ReLU()
  (dropout): Dropout(p=0.2)
  (batchnorm1): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (batchnorm2): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (batchnorm3): BatchNorm1d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
```

```
In [14]: model = MulticlassClassification(num_feature = NUM_FEATURES, num_class=NUM_CLASSES)
model.to(device)

criterion = nn.CrossEntropyLoss(weight=class_weights.to(device))
optimizer = optim.Adam(model.parameters(), lr=LEARNING_RATE)
print(model)
```

```
MulticlassClassification(
  (layer_1): Linear(in_features=2381, out_features=512, bias=True)
  (layer_2): Linear(in_features=512, out_features=128, bias=True)
  (layer_3): Linear(in_features=128, out_features=64, bias=True)
  (layer_out): Linear(in_features=64, out_features=11, bias=True)
  (relu): ReLU()
  (dropout): Dropout(p=0.2)
  (batchnorm1): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (batchnorm2): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (batchnorm3): BatchNorm1d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
```

```
In [11]: def multi_acc(y_pred, y_test):
y_pred_softmax = torch.log_softmax(y_pred, dim = 1)
_, y_pred_tags = torch.max(y_pred_softmax, dim = 1)

correct_pred = (y_pred_tags == y_test).float()
acc = correct_pred.sum() / len(correct_pred)

acc = torch.round(acc * 100)

return (acc)
```

```
In [12]: accuracy_stats = {
```

```

    'train': [],
    "val": []
}
loss_stats = {
    'train': [],
    "val": []
}

```

```

In [15]: for e in tqdm(range(1, EPOCHS+1)):

    # TRAINING
    train_epoch_loss = 0
    train_epoch_acc = 0
    model.train()
    for X_train_batch, y_train_batch in train_loader:
        X_train_batch, y_train_batch = X_train_batch.to(device), y_train_batch.to(device)
        optimizer.zero_grad()

        y_train_pred = model(X_train_batch)

        train_loss = criterion(y_train_pred, y_train_batch)
        train_acc = multi_acc(y_train_pred, y_train_batch)

        train_loss.backward()
        optimizer.step()

        train_epoch_loss += train_loss.item()
        train_epoch_acc += train_acc.item()

    # VALIDATION
    with torch.no_grad():

        val_epoch_loss = 0
        val_epoch_acc = 0

        model.eval()
        for X_val_batch, y_val_batch in val_loader:
            X_val_batch, y_val_batch = X_val_batch.to(device), y_val_batch.to(device)

            y_val_pred = model(X_val_batch)

            val_loss = criterion(y_val_pred, y_val_batch)
            val_acc = multi_acc(y_val_pred, y_val_batch)

            val_epoch_loss += val_loss.item()
            val_epoch_acc += val_acc.item()
        loss_stats['train'].append(train_epoch_loss/len(train_loader))
        loss_stats['val'].append(val_epoch_loss/len(val_loader))
        accuracy_stats['train'].append(train_epoch_acc/len(train_loader))
        accuracy_stats['val'].append(val_epoch_acc/len(val_loader))

    print(f'Epoch {e+0:03}: | Train Loss: {train_epoch_loss/len(train_loader):.5f} |

```

```

Epoch 001: | Train Loss: 0.41790 | Val Loss: 0.22116 | Train Acc: 88.328 | Val Acc: 92.053
Epoch 002: | Train Loss: 0.23859 | Val Loss: 0.20268 | Train Acc: 92.360 | Val Acc: 92.876
Epoch 003: | Train Loss: 0.19627 | Val Loss: 0.20098 | Train Acc: 93.507 | Val Acc: 92.924
Epoch 004: | Train Loss: 0.16484 | Val Loss: 0.19252 | Train Acc: 94.409 | Val Acc: 93.413
Epoch 005: | Train Loss: 0.14514 | Val Loss: 0.18975 | Train Acc: 94.925 | Val Acc: 93.484

```

Epoch 006:	Train Loss: 0.12980	Val Loss: 0.18957	Train Acc: 95.459	Val Acc: 93.974
Epoch 007:	Train Loss: 0.11753	Val Loss: 0.19869	Train Acc: 95.746	Val Acc: 94.081
Epoch 008:	Train Loss: 0.10754	Val Loss: 0.20689	Train Acc: 96.138	Val Acc: 93.890
Epoch 009:	Train Loss: 0.10033	Val Loss: 0.20554	Train Acc: 96.235	Val Acc: 94.368
Epoch 010:	Train Loss: 0.09180	Val Loss: 0.22072	Train Acc: 96.574	Val Acc: 93.950
Epoch 011:	Train Loss: 0.08742	Val Loss: 0.22180	Train Acc: 96.726	Val Acc: 94.069
Epoch 012:	Train Loss: 0.08105	Val Loss: 0.21773	Train Acc: 96.926	Val Acc: 94.308
Epoch 013:	Train Loss: 0.07668	Val Loss: 0.23866	Train Acc: 97.096	Val Acc: 94.021
Epoch 014:	Train Loss: 0.07139	Val Loss: 0.25575	Train Acc: 97.274	Val Acc: 94.165
Epoch 015:	Train Loss: 0.06980	Val Loss: 0.24474	Train Acc: 97.323	Val Acc: 94.332
Epoch 016:	Train Loss: 0.06365	Val Loss: 0.25575	Train Acc: 97.488	Val Acc: 94.189
Epoch 017:	Train Loss: 0.06236	Val Loss: 0.26237	Train Acc: 97.572	Val Acc: 94.117
Epoch 018:	Train Loss: 0.05824	Val Loss: 0.26963	Train Acc: 97.734	Val Acc: 94.224
Epoch 019:	Train Loss: 0.05700	Val Loss: 0.27907	Train Acc: 97.832	Val Acc: 93.902
Epoch 020:	Train Loss: 0.05291	Val Loss: 0.27230	Train Acc: 97.986	Val Acc: 94.177
Epoch 021:	Train Loss: 0.05105	Val Loss: 0.26992	Train Acc: 98.056	Val Acc: 94.344
Epoch 022:	Train Loss: 0.04678	Val Loss: 0.30489	Train Acc: 98.168	Val Acc: 94.272
Epoch 023:	Train Loss: 0.04538	Val Loss: 0.29419	Train Acc: 98.282	Val Acc: 94.379
Epoch 024:	Train Loss: 0.04311	Val Loss: 0.28303	Train Acc: 98.388	Val Acc: 94.284
Epoch 025:	Train Loss: 0.04053	Val Loss: 0.32409	Train Acc: 98.454	Val Acc: 94.558
Epoch 026:	Train Loss: 0.03950	Val Loss: 0.29567	Train Acc: 98.534	Val Acc: 94.463
Epoch 027:	Train Loss: 0.03631	Val Loss: 0.33275	Train Acc: 98.645	Val Acc: 93.735
Epoch 028:	Train Loss: 0.03532	Val Loss: 0.34068	Train Acc: 98.706	Val Acc: 94.248
Epoch 029:	Train Loss: 0.03483	Val Loss: 0.30766	Train Acc: 98.741	Val Acc: 94.379
Epoch 030:	Train Loss: 0.03177	Val Loss: 0.30858	Train Acc: 98.842	Val Acc: 94.714
Epoch 031:	Train Loss: 0.03059	Val Loss: 0.33766	Train Acc: 98.922	Val Acc: 94.177
Epoch 032:	Train Loss: 0.02927	Val Loss: 0.32903	Train Acc: 98.989	Val Acc: 94.153
Epoch 033:	Train Loss: 0.02803	Val Loss: 0.33556	Train Acc: 98.991	Val Acc: 94.379
Epoch 034:	Train Loss: 0.02742	Val Loss: 0.32791	Train Acc: 99.058	Val Acc: 94.153
Epoch 035:	Train Loss: 0.02613	Val Loss: 0.30905	Train Acc: 99.147	Val Acc: 94.666
Epoch 036:	Train Loss: 0.02455	Val Loss: 0.33800	Train Acc: 99.123	Val Acc: 94.200
Epoch 037:	Train Loss: 0.02455	Val Loss: 0.33508	Train Acc: 99.163	Val Acc: 94.511
Epoch 038:	Train Loss: 0.02352	Val Loss: 0.34711	Train Acc: 99.238	Val Acc: 94.391
Epoch 039:	Train Loss: 0.02235	Val Loss: 0.34183	Train Acc: 99.260	Val Acc: 94.320
Epoch 040:	Train Loss: 0.02218	Val Loss: 0.33779	Train Acc: 99.251	Val Acc: 94.320

```

4.332
Epoch 041: | Train Loss: 0.02232 | Val Loss: 0.36364 | Train Acc: 99.278 | Val Acc: 9
4.356
Epoch 042: | Train Loss: 0.02118 | Val Loss: 0.35037 | Train Acc: 99.327 | Val Acc: 9
4.368
Epoch 043: | Train Loss: 0.01996 | Val Loss: 0.34793 | Train Acc: 99.352 | Val Acc: 9
4.678
Epoch 044: | Train Loss: 0.01811 | Val Loss: 0.35122 | Train Acc: 99.388 | Val Acc: 9
4.785
Epoch 045: | Train Loss: 0.01849 | Val Loss: 0.35650 | Train Acc: 99.408 | Val Acc: 9
4.379
Epoch 046: | Train Loss: 0.01829 | Val Loss: 0.39512 | Train Acc: 99.383 | Val Acc: 9
4.308
Epoch 047: | Train Loss: 0.01701 | Val Loss: 0.40093 | Train Acc: 99.423 | Val Acc: 9
3.699
Epoch 048: | Train Loss: 0.01763 | Val Loss: 0.34270 | Train Acc: 99.420 | Val Acc: 9
4.499
Epoch 049: | Train Loss: 0.01704 | Val Loss: 0.34712 | Train Acc: 99.445 | Val Acc: 9
4.523
Epoch 050: | Train Loss: 0.01658 | Val Loss: 0.42581 | Train Acc: 99.480 | Val Acc: 9
4.177

```

```
In [26]: torch.save(model.state_dict(), os.path.join(sorel_dir, 'classification.pth'))
```

```
In [28]: os.path.join(sorel_dir, 'classification.pth')
```

```
Out[28]: '../Dataset/sorel/classification.pth'
```

```
In [30]: !ls ../Dataset/sorel/
```

```

classification.pth  sorel_label.csv          test-features.npz
db                  sorel_label_sum25.csv    validation-features.npz
meta.db             sorel_malware1.csv
sorel_data.npz      sorel_malware.csv

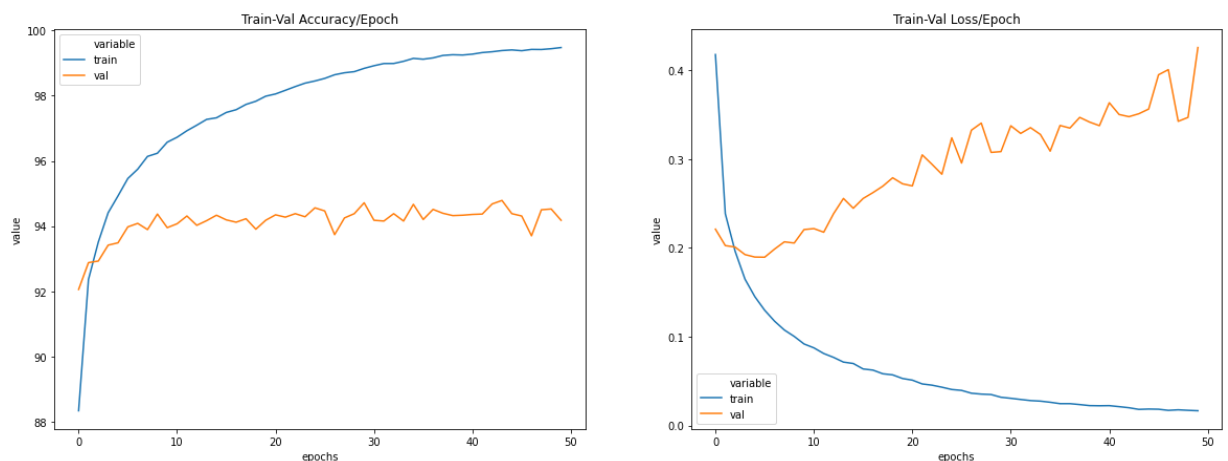
```

```

In [16]: # Create dataframes
train_val_acc_df = pd.DataFrame.from_dict(accuracy_stats).reset_index().melt(id_vars=
train_val_loss_df = pd.DataFrame.from_dict(loss_stats).reset_index().melt(id_vars=['
# Plot the dataframes
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(20,7))
sns.lineplot(data=train_val_acc_df, x="epochs", y="value", hue="variable", ax=axes[0])
sns.lineplot(data=train_val_loss_df, x="epochs", y="value", hue="variable", ax=axes[1])

```

```
Out[16]: Text(0.5, 1.0, 'Train-Val Loss/Epoch')
```



```

In [32]: y_pred_list = []
y_prob_list = []
with torch.no_grad():
    model.eval()
    for X_batch, _ in test_loader:

```

```

X_batch = X_batch.to(device)
y_test_pred = model(X_batch)
_, y_pred_tags = torch.max(y_test_pred, dim = 1)

y_pred_list.append(y_pred_tags.cpu().numpy())
y_prob_list.append(y_test_pred)

```

```
y_pred_list = [a.squeeze().tolist() for a in y_pred_list]
```

```
In [33]: y_prob_list[0], y_pred_list[0]
```

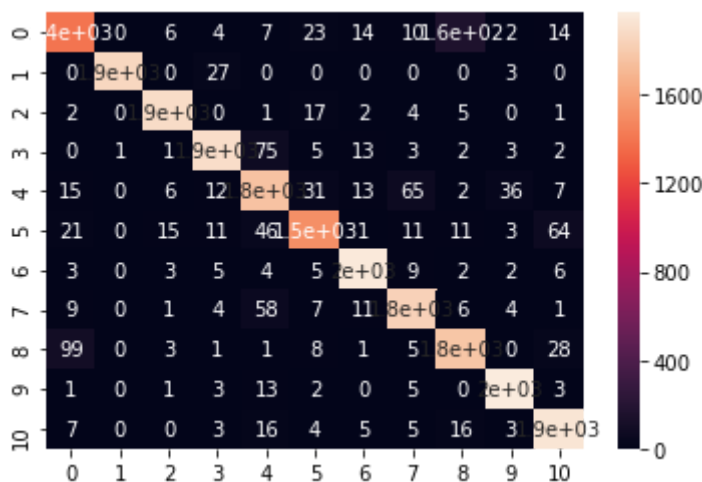
```
Out[33]: (tensor([[ -15.1259, -23.0870, -25.0988, -23.7968, -21.3749, -21.3992, -19.6963,
                   -17.8799,   5.0449, -22.6265, -18.0614]], device='cuda:0'),
          8)
```

```
In [16]: y_test.astype('long')
```

```
Out[16]: array([ 8,  6, 10, ...,  4,  4,  7])
```

```
In [17]: confusion_matrix_df = pd.DataFrame(confusion_matrix(y_test.astype('long'), y_pred_list),
                                             index=range(11), columns=range(11),
                                             sns.heatmap(confusion_matrix_df, annot=True))
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x2aed144500f0>
```



```
In [19]: confusion_matrix_df
```

```
Out[19]:
```

	0	1	2	3	4	5	6	7	8	9	10
0	1401	0	6	4	7	23	14	10	161	2	14
1	0	1866	0	27	0	0	0	0	0	3	0
2	2	0	1889	0	1	17	2	4	5	0	1
3	0	1	1	1895	75	5	13	3	2	3	2
4	15	0	6	12	1757	31	13	65	2	36	7
5	21	0	15	11	46	1510	1	11	11	3	64
6	3	0	3	5	4	5	1959	9	2	2	6
7	9	0	1	4	58	7	11	1841	6	4	1
8	99	0	3	1	1	8	1	5	1773	0	28
9	1	0	1	3	13	2	0	5	0	1972	3
10	7	0	0	3	16	4	5	5	16	3	1903

	0	1	2	3	4	5	6	7	8	9	10
10	7	0	0	3	16	4	5	5	16	3	1936

```
In [18]: print(classification_report(y_test.astype('long'), y_pred_list))
```

	precision	recall	f1-score	support
0	0.90	0.85	0.88	1642
1	1.00	0.98	0.99	1896
2	0.98	0.98	0.98	1921
3	0.96	0.95	0.96	2000
4	0.89	0.90	0.90	1944
5	0.94	0.89	0.91	1693
6	0.97	0.98	0.98	1998
7	0.94	0.95	0.94	1942
8	0.90	0.92	0.91	1919
9	0.97	0.99	0.98	2000
10	0.94	0.97	0.95	1995
accuracy			0.95	20950
macro avg	0.94	0.94	0.94	20950
weighted avg	0.95	0.95	0.94	20950

```
In [36]: ember2018 = '../Dataset/ember_zip/ember2018'
```

```
In [35]: EPOCHS = 50
BATCH_SIZE = 16
LEARNING_RATE = 0.0007
NUM_FEATURES = 2381
NUM_CLASSES = 2
```

```
In [64]: test_dataset = np.load(os.path.join(ember2018,'ember2018_test_data.npz'),allow_pickle=True)
X_train,y_train = test_dataset['arr_0'],test_dataset['arr_1']
```